


Decision Intelligence Runtime (DIR): An Architectural Pattern for Safe AI Execution

Bridging the gap between probabilistic reasoning and deterministic action

Author: Artur Huk

Repository: <https://github.com/huka81/decision-intelligence-runtime>

Last updated: 2026-01-05



0. Abstract

Most contemporary AI agent frameworks operate on a simple loop: observe, reason, and act. While effective for conversational tasks, this model fails in high-stakes environments where actions have financial or physical consequences. The core issue is architectural: Large Language Models (LLMs)

are probabilistic engines, yet they are often given direct control over deterministic interfaces (APIs, databases).

This paper introduces the **Decision Intelligence Runtime (DIR)**, an architectural pattern derived from two years of prototyping **Alvestor**, an autonomous algorithmic trading system. While born in the financial domain, DIR is not a trading-specific tool. It was designed to create a "Digital Investment Twin"—a system capable of understanding a user's strategy, executing transactions on their behalf, and rigorously documenting every decision. This pattern is equally applicable to any domain requiring auditable autonomy, from cloud infrastructure management to supply chain logistics.

DIR applies principles from distributed systems orchestration (sagas, idempotency) and security (policy enforcement points) to the domain of AI agents. It proposes a strict separation of concerns where agents are responsible for **Reasoning** (proposing strategies) and a deterministic runtime is responsible for **Execution** (validating and applying those strategies).

By decoupling intent from action, DIR solves common stability issues such as race conditions, hallucinations in function calls, and execution of stale decisions. This document outlines the pattern's core components, including the DecisionFlow ID (an adaptation of distributed tracing for reasoning chains) and the Decision Integrity Module, offering a blueprint for moving agents from experimental scripts to reliable production systems.

1. Motivation: Why Agents Need a Runtime

Over the last two years, I built and operated Alvestor, an autonomous system designed to manage a virtual trading portfolio. The initial implementation followed the standard "agentic" pattern popular in the industry: an LLM loop that analyzed market data and directly called broker APIs.

The results were technically impressive but operationally terrifying.

The system was capable of sophisticated reasoning but lacked execution discipline. It would occasionally attempt to sell positions it no longer held because of a state update lag. It would sometimes "hallucinate" a trade retry loop, ignoring API rate limits. Most critically, it treated time as an abstract concept; a "buy" decision made based on a price from 10 seconds ago would be executed 30 seconds later, often incurring slippage that invalidated the original strategy.

These were not failures of intelligence. They were failures of architecture.

1.1 The "Probabilistic to Deterministic" Gap

The fundamental problem in modern agent design is the collapse of two distinct concerns into a single loop:

1. **Reasoning (Probabilistic):** The agent interpreting context. This is messy, creative, and non-deterministic.
2. **Execution (Deterministic):** The system changing state (e.g., sending money, updating a record). This requires strict guarantees.

In standard software engineering, we solve similar problems using patterns like **CQRS** (Command Query Responsibility Segregation). We separate the *intent* to change data from the *process* of changing it. Yet, in most AI frameworks, we allow the probabilistic model to write directly to the "database" of the real world.

When reasoning and execution are interleaved, non-determinism leaks into operational behavior. Safety mechanisms become prompts ("Please do not trade if volatility is high") rather than hard constraints. As any security engineer knows, prompts are not permissions.

1.2 Moving Beyond "Prompt Engineering" to "System Engineering"

To stabilize Alvestor, I had to stop treating it as a chatbot and start treating it as a distributed system. I realized that reliable agents require the same infrastructure we use for microservices, adapted for the unpredictability of LLMs:

- **Orchestration:** Just as we use tools like Temporal or Cadence to manage long-running workflows, agents need a runtime to manage the lifecycle of a decision.
- **Idempotency:** Agents will repeat themselves. The system must recognize duplicate intents and prevent duplicate side effects (e.g., executing the same trade twice).
- **Traceability:** In a microservice, we use a **Trace ID** (OpenTelemetry) to follow a request. In an agent system, we need to trace the *reasoning chain* that led to an action. This concept evolved into what I call the **DecisionFlow ID (DFID)**.
- **Time-to-Live (TTL):** Data expires. An agent's intent must have a strict validity window. If the runtime cannot execute the decision within that window, it must be discarded, not delayed.

1.3 From Experiment to Pattern

The Decision Intelligence Runtime (DIR) is not a software product. It is a set of architectural constraints and patterns designed to make AI systems auditable and safe.

It shifts the design philosophy from **Agent-Centric** (how smart is the model?) to **System-Centric** (how robust is the execution?).

- Agents answer: "What should we do and why?"
- The Runtime answers: "Is this action allowed, valid, and safe to execute right now?"

A note on terminology: Throughout this document, 'ROA' refers to **Responsibility-Oriented Agents**, a pattern for bounding AI autonomy, unrelated to the older Resource-Oriented Architecture definition.

The following sections define the components of this runtime, illustrating how to wrap "fuzzy" agent logic in a "hard" engineering shell.

2. Defining the Scope: The Runtime as Middleware


In the early iterations of Alvestor, the agent code was monolithic. The same Python script was responsible for parsing news, deciding on a strategy, and sending HTTP requests to the broker. This tightly coupled design meant that a bug in the reasoning logic (e.g., a loop caused by a misunderstood prompt) resulted in direct operational hazards.

To fix this, I adopted a pattern familiar to OS developers: **Kernel Space vs. User Space separation**.

In this architecture, the **Decision Intelligence Runtime (DIR)** acts as the Kernel. It manages resources, enforcing permissions and time constraints. The Agents operate in User Space; they can request actions, but they cannot execute them directly.

2.1 The Architectural Layering

DIR sits strictly between the probabilistic agents and the deterministic infrastructure.



Its responsibilities are scoped to:

1. **Orchestration:** Receiving unvalidated proposals from agents and processing them through a deterministic pipeline.
2. **Validation:** Functioning as a **Policy Enforcement Point (PEP)**^[1], similar to OPA (Open Policy Agent)^[2] in cloud-native security.
3. **Translation:** Converting "soft" agent intents (policies) into "hard" execution commands (API calls) with idempotency guarantees.
4. **Feedback:** Closing the loop by returning `ValidationFeedback` events to the Agent. When a policy is rejected (e.g., `RISK_LIMIT_EXCEEDED`), the Runtime must inform the Agent *why*, so it can attempt self-correction in the next cycle.

2.2 What It Is Not (Scope Exclusion)

Crucially, the Runtime is **not** an Agent. It contains no LLMs and performs no semantic reasoning. If the system needs to "think" or "interpret," that belongs to the agent layer. The Runtime is purely a state machine designed to govern the side effects of that thinking.

2.3 The Agent Registry: System of Record for Responsibilities

In a static system, hard-coding agent permissions works. In Alvestor, as specialized agents (e.g., "Momentum Trader", "Hedge Manager") were added and removed dynamically, hard-coding failed.

DIR introduces an **Agent Registry**-a service discovery^[3] mechanism for intelligence.

- **Registration:** On startup, an agent registers its `Manifest`: its ID, its subscribed inputs (Context), and its authorized outputs (Policy Types).
- **Capability Contract:** The Registry acts as the source of truth for ROA constraints. When the Validation Layer asks "Can Agent X trade Asset Y?", it queries the Registry, not the Agent. This prevents agents from self-granting permissions via prompt injection.

Conceptual Decomposition (Single Service, Multiple Authorities)

*Although implemented as a single service, the Registry fulfills multiple **conceptual authorities**:

- Identity & Capability Authority
- Schema Authority
- Reservation / Lock Authority
- Lifecycle Authority

This decomposition is conceptual, not necessarily physical, and exists to prevent semantic overloading of responsibility.*

Beyond capability tracking, the Agent Registry facilitates **Resource Locking and Reservation**. In environments where multiple agents (e.g., concurrent PositionAgents) operate on a shared finite resource—such as a single capital pool or a limited API throughput—the Registry acts as a synchronization point. It allows the Runtime to grant temporary 'Reservation Locks' to a DecisionFlow.

Operational Resilience (Addressing SPOF):

The Registry is a critical component. To prevent it from becoming a Single Point of Failure (SPOF), the Runtime implements **Local Manifest Caching**.

- **Cache Strategy:** The Runtime caches Agent Manifests locally with a short TTL (e.g., 60 seconds).
- **Degraded Mode:** If the Registry is unreachable, the Runtime continues to serve known agents using cached definitions. New agent registrations or schema updates are rejected until connectivity is restored.

Note on Schema Evolution:

This dynamism requires that Agents do not "memorize" the policy schema indefinitely. Instead, the Agent Registry serves the current version of the JSON schema dynamically during the Context compilation step. This ensures that even as capabilities evolve, the Agent always reasons against a valid, up-to-date interface contract.


Strict Versioning (Avoiding "Contract Hell"):

In distributed agent systems, mismatched expectations lead to failures. The Agent Registry mandates **Semantic Versioning (SemVer)** alignment. An agent initialized with `v1.2` capability manifests must negotiate with a Runtime supporting `v1.x` schemas. If a disconnect is detected, the Runtime rejects the agent's registration during the handshake, preventing runtime parsing errors in production.

Registry Updates and Flow Binding:

Registry updates are versioned. A `DecisionFlow` is bound to the Registry snapshot version active at `CREATED` state. Any mid-flow authority revocation (e.g., security kill-switch) triggers an immediate `ABORT` on the next JIT check, ensuring no policy executes under revoked permissions.

The Capabilities Handshake - Startup Sequence



3. System Invariants

Instead of a "manifesto," DIR relies on a set of architectural invariants. These are the constraints that must hold true for the system to be considered safe, regardless of how "creative" the LLM becomes.

3.1 Invariant 1: Deterministic State Transitions

While the inputs to the system (User Space reasoning) are probabilistic, the transition from a **Validated Policy** to an **Execution** (Kernel Space) must be deterministic.

- **User Space (Probabilistic):** Agents, LLMs, prompts.

- **Kernel Space (Deterministic):** Validation logic, state machines, API calls.

Given the same Policy Proposal, the same Context Snapshot, and the same Time, the Runtime must always produce the exact same Validation Result. This requires that **Hard Gates** (blocking validation) be implemented in standard code (Python/Go/Rust) or a policy engine (Rego). Probabilistic validation (LLM-based checks) must remain in the User Space or serve as non-blocking observers, unless explicitly configured otherwise (see Sec 6.3).

3.2 Invariant 2: The "Reasoning-Execution" Wall

This is an adaptation of the **Command Query Responsibility Segregation (CQRS)**^[4] pattern.

- **Agents (Write Model via Proposal):** Agents perform the reasoning and emit a `PolicyProposal`. This is equivalent to a Command in CQRS, but with a critical distinction: it is *tentative*.
- **Runtime (Execution):** The Runtime validates the proposal. Only if validation passes does it trigger a side effect.
- **Constraint:** No agent is ever permitted to hold API keys or database write credentials. Agents only have permission to "submit proposals" to the Runtime's internal bus.

3.3 Invariant 3: Execution Parametrization (Constraints over Deadlines)

In distributed systems, we often use "Deadline Propagation." However, relying solely on a hard TTL (Time-To-Live) for AI decisions is brittle. If an LLM takes 8 seconds to "think" and the queue takes 3 seconds, a 10s TTL rejects valid decisions.

DIR replaces robust "race-against-the-clock" logic with **Execution Parametrization**.

- **Logic:** The Agent does not propose: *"Buy at the price I saw in the snapshot (\$100)."*
- **Constraint:** The Agent proposes: *"Buy with a limit of \$102 (Acceptable Slippage: 2%)."*
- **Mechanism:** The Runtime checks the **Execution Constraints** at the exact moment of execution. This effectively decouples "slow" reasoning time from "fast" execution time, ensuring that latency does not invalidate the strategy, provided the market conditions remain within the Agent's defined bounds.

3.4 Invariant 4: Auditability by Correlation

To debug a distributed system, we use Trace IDs. To debug an AI system, we need to trace the *causality*.

Every artifact in the system (from the initial observation to the final API response) is tagged with a **DecisionFlow ID (DFID)**. This allows us to reconstruct the entire narrative: *Context -> Prompt ->*

Reasoning -> Proposal -> Validation -> Execution. Without this correlation, explaining *why* the system lost money (or crashed) is impossible.

4. DecisionFlow: Distributed Tracing for Reasoning

In microservices, we use **Distributed Tracing**^[5] (e.g., OpenTelemetry) to follow a request across service boundaries. We assign a `TraceID` at the ingress and propagate it everywhere.

In AI agents, the complexity lies not just in *where* the request went, but *how* the decision was formed. Standard application logs show "Database Updated," but they don't show the prompt, the context snapshot, or the LLM's rationale that led to that update.


4.1 The DFID (DecisionFlow ID)

To solve this in Alvestor, I introduced the **DecisionFlow ID (DFID)**. Conceptually, this is a Correlation ID, but it spans a wider scope than a typical HTTP request. A DFID acts as a container for the entire lifecycle of a single intent.

All operations—observations, prompts, reasonings, and execution results—are persisted in a database, tagged with this ID. In an event-driven implementation, this ID propagates through the Event Bus, allowing subscribers (like an Audit Service) to reconstruct the full causality chain.

It binds together:

1. **The Trigger:** The market event or timer that woke the agent up.
2. **The Context Snapshot:** A hash or link to the exact data the agent "saw" (crucial for replayability).
3. **The Reasoning:** The raw LLM output explaining *why*.
4. **The Policy Proposal:** The structured JSON intent.
5. **The Validation Outcome:** Why the runtime accepted or rejected it.
6. **The Execution Result:** The final side effect (e.g., transaction ID).



4.2 Hierarchical DecisionFlows


In complex domains, decisions are rarely atomic. A strategic decision ("Reduce Tech Exposure") often spawns multiple tactical decisions ("Sell AAPL", "Buy Put Options").

DIR supports **Parent-Child relationships** between DFIDs.

- **Parent Flow:** Represents the high-level intent (Strategy).
- **Child Flows:** Represent the granular actions (Execution).

This hierarchy allows for precise auditing. If a trade fails, we can trace it back not just to the specific tactical agent, but to the parent strategic mandate that authorized it.

The "Umbrella" Pattern - One Strategy, Many Actions



4.3 Lifecycle Management

Unlike a stateless HTTP request, a DecisionFlow is a stateful entity. It follows a strict lifecycle managed by the Runtime:

- **CREATED:** The flow is initialized.
- **ACTIVE:** The agent is reasoning or a proposal is being validated.
- **CLOSED:** Execution completed successfully.
- **ABORTED:** The flow was terminated due to validation failure, timeout (TTL), or error.

5. The Interface: Policies as Contracts

The most common mistake in agent development is allowing the LLM to output free text or directly generate code (e.g., SQL). This makes the system unpredictable and hard to parse.

In DIR, the interface between the Agent (Reasoning) and the Runtime (Execution) is strictly defined by a schema. We call this the **Policy**.

5.1 Agents Propose, Runtimes Dispose

This follows the **Declarative API** pattern^[6] (similar to Kubernetes manifests).

- The Agent does not say: *"Call the API to buy Apple stock."* (Imperative)
- The Agent emits a Policy Object: *"I desire a state where we own 10 shares of AAPL, given current price X."* (Declarative)

The Runtime then evaluates if this desired state is permissible.

5.2 The "Explain vs. Policy" Separation

In practice, I found that LLMs perform better when allowed to "think out loud" before committing to a format. Therefore, the Agent's output is split into two distinct channels:

1. **Explain (Unstructured):** A natural language narrative for human auditors. **It provides human context.**
2. **Policy (Structured):** A strict JSON/Pydantic object containing the specific parameters for the action. **It provides machine-interpretable intent.**

The Runtime validates *only* the Policy. The Explanation is treated as metadata (comments). This separation prevents the system from mistaking a narrative justification for an executable instruction.

Intent vs. Execution Tactics

It is important to clarify that a Policy Proposal is not a raw market order (e.g., "Buy at market NOW"). Given the latency of LLM inference, such an approach would suffer from inevitable slippage. Instead, ROA Agents emit **Strategic Intent**s with explicit **Execution Constraints**.

The Agent defines the *boundary conditions* (e.g., "Buy up to \$102"), and the Runtime executes within those bounds. This separates the "thinking time" from the "market timestamp."

Example Structure (JSON):

```

{
  "dfid": "550e8400-e29b-41d4-a716-446655440000",
  "agent_id": "risk_manager_v1",
  "policy_kind": "ADJUST_POSITION",
  "execution_constraints": {
    "max_slippage_bps": 50,
    "validity_window_sec": 30,
    "requires_market_state": "OPEN"
  },
  "params": {
    "symbol": "BTC-USD",
    "action": "REDUCE",
    "quantity": 0.5
  },
  "context_ref": "snapshot_hash_x9823"
}

```

5.3 Claims vs. Facts

A Policy Proposal is treated as a **Claim** (an untrusted assertion), not a **Fact**.

The Agent *claims* that selling Bitcoin is the right move. It becomes a *fact* (an executed event) only after the Runtime validates the signature, the permissions, and the market state. This distinction prevents the "authority bias" where we implicitly trust the AI just because it produced an output.

6. Validation: The Deterministic Gate (DIM)

In the "User Space vs. Kernel Space" analogy, the **Decision Integrity Module (DIM)** is the kernel's access control list. It is the gatekeeper that determines whether a User Space proposal is allowed to touch the infrastructure.

6.1 Deterministic Enforcement, Not Probabilistic Guessing

A common anti-pattern in agent systems is "LLM-based validation"-asking a second LLM to critique the first one. While useful for improving reasoning quality, this is insufficient for safety.

In DIR, the validation layer is strictly **deterministic**. It is implemented in code (e.g., Python, Go, or Rego policies), not prompts. Given inputs (Policy, Context, Time), the output must always be

ACCEPT OR REJECT .

6.2 The Validation Pipeline

The pipeline functions as a **Policy Enforcement Point (PEP)**. It evaluates proposals against three layers of constraints:

1. **Schema & Integrity:** Does the JSON match the versioned schema?
2. **Authority (RBAC):** Is this agent authorized in the *Agent Registry* to execute this Policy Kind?
3. **State Consistency (Optimistic Concurrency)**^[7]: Does the `context_hash` in the proposal match the current system state? If slippage occurred, reject with `STALE_CONTEXT` .
4. **Resource Availability (Semantic Locking):** To prevent "Horizontal Resource Contention" (where two agents compete for the same cash/inventory), the DIM places a temporary lock or reservation on the required assets during the validation phase. If Agent A has reserved the last unit of capital, Agent B's simultaneous request is rejected with `INSUFFICIENT_LIQUIDITY` , preventing race conditions.
 - **Linear Lock Acquisition:** To prevent deadlocks in multi-resource requests, resources **MUST** be requested in alphabetical order of their Global Resource IDs. Failure of the Agent to adhere to this sorting order in the Policy Proposal results in immediate rejection by the DIM. A mandatory `LockTimeout` (e.g., 5s) ensures that stalled flows are `ABORTED` with `RESOURCE_CONTENTION_TIMEOUT` .
5. **Mission Invariant Check:** The DIM **MUST** verify that the Policy Proposal contains a `mission_context_hash` . The Runtime compares this against the registered Agent Mission. If the agent's reasoning context has drifted from its assigned mission, the DIM rejects the proposal with `MISSION_DISSONANCE` .

The Runtime does not interpret mission semantics.

*It validates **contractual alignment**, not semantic intent.*

The `mission_context_hash` represents an immutable contract snapshot, not a philosophical goal.

Drift & Explanation-Intent Divergence

Semantic Alignment is not a security mechanism.

It does not prevent malicious or incorrect actions.

*It exists to preserve **operator trust and audit clarity**, not system safety.*

The pipeline incorporates an **Intent Retry Governor** to mitigate the risk of 'Feedback Poisoning.' When the Runtime returns `ValidationFeedback` to an agent following a rejection (e.g., a risk limit violation), the agent is permitted a strictly limited number of attempts (`Maximum Intent Retries`, typically 3) to correct its proposal within the same `DecisionFlow`. If the agent continues to produce non-compliant policies after these attempts, the Runtime forcibly terminates the flow with a

REASONING_EXHAUSTION status. This protects the system from infinite reasoning loops and prevents a hallucinating model from draining the Token Budget through unproductive attempts to bypass deterministic guards.

6.3 Semantic Alignment Check (The "Liar" Detection)

A subtle failure mode in LLMs is "proxy gaming," where the model's narrative ("I am reducing risk") contradicts its structured policy ({ "action": "BUY_LEVERAGE" }).

To counter this, DIR supports an optional **Semantic Alignment Check**.

Hard Gates vs. Soft Guards

To preserve the determinism of the runtime (Invariant 1), we distinguish between:

1. **Hard Gates (Deterministic):** Rego policies, schema validation, RBAC code, and arithmetic checks. These are **blocking**.
2. **Soft Guards (Probabilistic):** LLM-based semantic checks. These operate strictly as **Auditors**.

Default Behavior: Audit Only

By default, if a Soft Guard detects a mismatch (SEMANTIC_MISMATCH), it **does not** block execution. Instead, it:

- Flags the resulting DecisionFlow as NEEDS_REVIEW (Post-Execution Audit).
- Triggers an asynchronous alert to the operator.
- Allows the transaction to proceed if all Hard Gates are satisfied.

Strict Mode (Optional Exception)

For systems where safety is paramount over availability, an architect may enable `strict_semantic_blocking: true`.

- **Effect:** A Semantic Mismatch triggers an immediate ABORT.
- **Warning:** This configuration **violates Invariant 1 (Determinism)**. A non-deterministic model update could cause previously valid transactions to fail. This mode is recommended only for low-throughput, high-risk environments where false positives are an acceptable cost.

6.4 Time as a Hard Constraint

The Runtime enforces the **Decision Validity Window**. If `current_time > policy.valid_until`, the proposal is rejected immediately. This prevents the "queued command" problem where a backlog of old decisions suddenly executes hours later.

6.5 Just-In-Time State Re-verification (The Anti-TOCTOU Check)

Validation occurs *before* execution, but in high-concurrency systems, the state may change while the request is in flight. This creates a **Time-of-Check to Time-of-Use (TOCTOU)** vulnerability.

To mitigate this, the Execution Engine enforces a **Just-In-Time (JIT) State Check**.

- **Mechanism:** Immediately before dispatching the `ExecutionIntent` (after locking limits but before the network call), the Runtime performs a lightweight assertion against the live `Authoritative Context`.
- **Assertion:** Verifies that critical invariants (e.g., `current_price` is roughly equal to `snapshot_price`, `balance` $\geq$ `required_amount`, `record_version` matches).
 - **Drift Tolerance:** State drift tolerance **MUST** be explicitly defined in the `ExecutionConstraints`. Defaulting to 'absolute' matching for balances and 'threshold-based' (e.g., $<0.1\%$) for telemetry.
- **Action:** If the state has drifted beyond the allowed tolerance or if the context age exceeds a hard threshold (e.g., `current_state_age` $> 500\text{ms}$), the Runtime aborts with `STATE_DRIFT_DETECTED` and forces the Agent to re-reason.

Architectural Note: This introduces a performance penalty (an extra read operation). DIR accepts this cost ("Safety over Speed") to guarantee that no decision executes against a phantom reality.

7. Execution: Idempotency and Side Effects

Once a policy is accepted, the system must cross the "Rubicon" into the real world. This is where we encounter the messy reality of external APIs: timeouts, network partitions, and 500 errors.

7.1 The Golden Rule: No Side Effects Without Intent

DIR enforces a strict security boundary:

No side effect may occur without an explicit Execution Intent.

Agents (User Space) never hold API keys or database credentials. They cannot open sockets. They can only submit proposals to the Runtime. The Runtime, after validation, transforms the `PolicyProposal` into an `ExecutionIntent`. This object is the only artifact in the system authorized to trigger external IO.

7.2 Idempotency: The "Double-Spend" Protection


LLMs can get stuck in loops, and network retries can deliver duplicate messages. To protect against this, DIR assigns a deterministic **Idempotency Key**^[8] to every Execution Intent.

The Key Formula:

$\text{IdempotencyKey} = \text{SHA256}(\text{DFID} + \text{Step_ID} + \text{Canonical_Params})$

- **DFID:** The trace ID of the reasoning chain.
- **Step_ID:** For multi-step sequences.
- **Canonical_Params:** A sorted string of the action parameters.

The `Attempt_Number` is logged for observability but **MUST NOT** be part of the key. This ensures that retries of the same intent resolve to the same side effect. If the Runtime sees a duplicate key, it returns the *cached result* of the previous execution rather than triggering the API again.



Atomicity and Parent-Child Sagas

DIR treats every `ExecutionIntent` as **Atomic**. The Runtime does not natively support "multi-step transactions" within a single intent because they are typically non-deterministic during failure.

Parent-Child Sagas

For complex workflows (e.g., "Sell Asset A to fund purchase of Asset B"), DIR utilizes a **Parent-Child DecisionFlow** pattern:

1. **Parent Agent (Saga Manager):** Maintains the state of the complex transaction. It emits a Policy to spawn a Child Flow for Step 1.
2. **Child Agent (Executor):** Receives the mandate, acts atomically (e.g., "Sell A"), and reports success/failure to the Parent.
3. **Failure Handling:** If Step 1 succeeds but Step 2 fails, the *Runtime* does not guess how to rollback. Instead, it reports the failure to the Parent Agent. The Parent Agent then reasons about the partial state and emits a new Policy: **Compensation Action** (e.g., "Re-buy Asset A" or "Alert Human").

This keeps the Runtime "dumb" and deterministic, while moving the complex recovery logic back to the entity capable of reasoning: the Agent

Not all external APIs are transactional/atomic. A Policy might require executing a sequence of dependent actions (e.g., "Sell Asset A to fund purchase of Asset B").

DIR rejects the "all-or-nothing" fantasy.

- **State: PARTIAL_SUCCESS_DIRTY:** If a 3-step policy fails at step 2, the DFID is not simply "failed." It is tagged as `DIRTY`.
- **Compensation:** This triggers a **Saga Compensation** workflow^[9]. Unlike a simple retry, this logic attempts to undo Step 1 or flag the anomaly for human resolution. Dirty states freeze the Agent instance in `MAINTENANCE_MODE` until a Compensation Policy is executed or human intervention clears the lock.

8. Context Management: Compilation over Conversation

In standard chatbots, "context" is simply the chat history. In an autonomous system like Alvestor, treating the entire event log as context is dangerous. It leads to **Context Window Overflow** and "distraction."

DIR treats context not as a log, but as a **Compiled Artifact**.

8.1 The 4 Layers of Context

To organize information effectively, the **Context Store** is structured into four distinct layers, each with different persistence and retrieval properties:

1. **Session (Ephemeral):** The append-only record of the current DecisionFlow (observations, proposals, validation results). It resets when the flow closes.
2. **State (Authoritative):** The current, trusted view of the world (e.g., wallet balance, open positions). This is often a read-replica of the external system state.
3. **Memory (Long-Lived):** Curated insights that persist across sessions (e.g., "Strategy A failed in high volatility").
4. **Artifacts (Reference):** Large blobs referenced by pointers (e.g., PDF reports, datasets) that are too large to fit in the prompt but available for tool-use retrieval.


8.2 Context Compilation Pipeline

Agents do not query the database directly. Instead, the Runtime executes a deterministic **Context Compilation** step before invoking the agent. This is analogous to the **Retrieval-Augmented Generation (RAG)**^[10] pattern, but strictly structured.

Context-Schema Validation:

The `WorkingContext` is version-stamped. Before invocation, the Agent Registry validates that the Agent's `ReasoningEngineVersion` is compatible with the `ContextSchemaVersion`. This prevents agents from interpreting malformed or outdated data snapshots after a system update.

The Compiler filters noise, enforcing a "Need-to-Know" policy. To prevent context window overflow, strict limits (e.g., "last X news items", "positions opened in last 24h") are enforced at the query level, ensuring the agent sees only the most relevant, recent slice of reality.



- **Inputs:** Raw Event Log, Market State, Static Rules.
- **Process:** Filter by Time -> Filter by Relevance (RAG) -> Format (JSON/Text).
- **Output:** `WorkingContext` object passed to the LLM.

8.3 Implementation Logic

The compilation process handles the **External State Paradox** (where the world changes while the agent thinks). It snapshots the state *at the moment of invocation*.

Context Compiler

```
def compile_working_context(agent_id: str, dfid: str) -> dict:
    # 1. Fetch Authoritative State (Snapshot)
    current_state = state_store.get_snapshot(timestamp=now())

    # 2. Retrieve relevant history (Session)
    session_events = event_log.query(dfid=dfid, limit=10)

    # 3. Retrieve static instructions (Memory)
    mission = memory_store.get_mission(agent_id)

    # 4. Assemble Immutable Context Object
    return {
        "snapshot_id": current_state.hash, # CRITICAL: This ID is verified by DIM (Sec 6.2) to
        "market_data": current_state.data,
        "recent_history": session_events,
        "mission": mission
    }
```

9. Human-in-the-Loop: Governance by Exception

The goal of autonomy is to reduce human toil, yet many agent frameworks default to a "Mother-May-I" pattern where a human must approve every single step. This leads to **Alert Fatigue**.

DIR adopts a **Governance by Exception** model. The system acts autonomously, it initiates an escalation routine.

Hierarchical Escalation (Agent -> Agent)

Before alerting a human, the system attempts **Hierarchical Escalation**:


1. **Peer/Supervisor Review:** A `PositionAgent` blocked by a risk limit can escalate the decision to its parent `StrategyAgent`.
2. **Scope Expansion:** The `StrategyAgent` has a broader context and higher authority limits. It may approve the action (override), modify the mandate, or reject it.
3. **Human Alert:** Only if the `StrategyAgent` also fails to resolve the issue (or lacks authority), does the system trigger a "Human-in-the-Loop" alert. This reduces "Alert Fatigue" for operators.

When the flow finally transitions to `ESCALATED` (Human required), the Runtime pauses execution and awaits external input

9.1 Escalation as a System State

Escalation is not an error; it is a valid state transition in the DecisionFlow.

When the Runtime encounters a situation it cannot resolve deterministically (e.g., ambiguity, risk limit violation, or repeated API failures), it transitions the flow to `ESCALATED`.



- *Nodes:* CREATED -> ACTIVE -> (Validation) -> [ACCEPTED | REJECTED | ESCALATED].
- *Transitions:*
- ACCEPTED -> EXECUTION -> CLOSED.
- ESCALATED -> (Human Review) -> [RESUME | ABORT].

9.2 Triggers for Intervention

In Alvestor, I defined clear criteria for when the machine must wake the human:

1. **Authority Violation:** Agent attempts to trade >\$1000 (Hard Limit).
2. **Model Uncertainty:** Agent produces a Policy with confidence < 0.7.
3. **Operational Failure:** Broker API returns 5xx error more than 3 times.
4. **Silence Watchdog:** No decision produced for >30 minutes during market hours.

9.3 Escalation Throttling & Cost Control

A failing agent can generate hundreds of escalation requests per minute, essentially DDoS-ing the human operator or the wallet.

The Escalation Budget

Each agent has a token bucket for escalations (e.g., 3 per hour). If the budget is exhausted, the agent is automatically demoted to a `PASSIVE` (Read-Only) state, and the DecisionFlow is silently aborted.

Computation Budget (Token Cap per DFID)

DIR extends budget control to compute costs. Each DecisionFlow is assigned a hard limit (e.g., \$0.50 or 10k tokens).

- **Mechanism:** Middleware tracks token accumulation across the reasoning chain.

- **Enforcement:** If an agent gets stuck in a "Chain-of-Thought" loop and fails to emit a Policy Proposal before the budget is drained, the Runtime aborts the operation (Timeout and Reject). This prevents "financial DDoS" from a buggy, rambling model.

Protection against "Feedback Poisoning" (Maximum Intent Retries)

A rejection by the DIM often triggers an agent retry loop. There is a specific risk that an agent, trying to bypass a safety check (e.g., Risk Limit), begins to "hallucinate compliance" or argue with the validator without changing the underlying parameters.

- **Mechanism:** The Runtime enforces `Maximum Intent Retries` (e.g., 3 attempts per DFID).
- **Enforcement:** If validation fails 3 times, the `DecisionFlow` is strictly `ABORTED` (not escalated). The system assumes the agent is caught in a "delusion loop" or "feedback poisoning" cycle and must be reset to prevent resource exhaustion.

9.4 The Human Interface

When a flow is escalated, the human acts as a **Super-User**.

The Runtime presents the `WorkingContext`, the `PolicyProposal`, and the `Reason` for escalation. The human then issues a binding decision: **OVERRIDE**, **MODIFY**, or **ABORT**.

Defense Against "Rubber Stamping":

A critical risk in Human-in-the-Loop systems is operator fatigue, leading to reflexive approval. The UI must never offer a simple "Approve" button next to a raw JSON block.

Requirement: Impact Categorization & Context

While a full "State Diff" simulation is the gold standard, it is often technically prohibitive. As a practical invariant, the Runtime must assign an **Impact Category** to every escalation:

- **LOW_IMPACT (Informational/Reversible):** e.g., "Change Logging Level", "Update Watchlist".
- **HIGH_IMPACT (Financial/Irreversible):** e.g., "Transfer Funds", "Execute Trade", "Delete Record".

The UI must visualize these categories (e.g., Red borders for `HIGH_IMPACT`) to disrupt "click-through" behavior. The operator approves the *risk category*, not just the JSON syntax.

[Escalation Event]

```
{
  "event_type": "ESCALATION_REQUIRED",
  "dfid": "trace_123_abc",
  "reason": "RISK_LIMIT_EXCEEDED",
  "details": {
    "proposed_value": 1500,
    "limit": 1000
  },
  "status": "AWAITING_HUMAN_INPUT"
}
```

10. Implementation Topologies: From Monolith to Mesh

A runtime pattern is only useful if it scales. While the logic of DIR (Validation, Execution, Context) is universal, the deployment topology changes as the system grows.

DIR does not mandate a specific architecture. It acts as a **substrate** that can host various agent interactions.

10.1 The Single-Process MVP (The "Alvestor" Model)

For most prototypes and early-stage production systems (like the current version of Alvestor), DIR runs effectively as a modular monolith.

- **Infrastructure:** Single Python process or Container.
- **Communication:** In-memory queues (e.g., Python `asyncio.Queue`).
- **State:** Local SQLite or PostgreSQL.
- **Pros:** Easy to debug, zero network latency between Agent and Runtime.
- **Cons:** Single point of failure, hard to scale processing power.


10.2 The Distributed Event-Driven Architecture (EDA)

For enterprise scale, DIR maps naturally to an **Event-Driven Architecture**^[11].

- **Infrastructure:** Microservices (Agents are services, Runtime is a service).
- **Communication:** Message Bus (e.g., Kafka, RabbitMQ, NATS).
- **State:** Distributed Key-Value Store (Redis) + Time Series DB.

- **Pattern:** The **Event-Oriented Agent Mesh**. Agents emit `PolicyProposal` events to a topic; the Runtime consumes them, validates, and emits `ExecutionIntent` events.

Topology Comparison - Monolith vs Distributed Architecture



- *Left:* Agents inside the same box as Runtime.
- *Right:* Agents and Runtime connected by a Message Bus.

This flexibility allows teams to start small (Monolith) and refactor to microservices (Mesh) without changing the core decision logic or validation rules.

11. Case Study: Lessons from Alvestor

The principles in this document were not derived from academic theory. They were reverse-engineered from failures encountered while building **Alvestor**, an autonomous trading prototype

running continuously for the past year.

11.1 The "Slippage" Incident (Why we need Time Validity)

Early versions of Alvestor had no concept of Decision Validity Windows. On one occasion, the system experienced a latency spike due to API rate limiting. An agent calculated a trade based on a price of 100. *The execution happened 45 seconds later when the price was 102.* The system executed the trade, locking in an immediate loss.

- *Lesson:* A decision is only good for the moment it was made. This led to the implementation of strict **TTL (Time-To-Live)** on all proposals.

11.2 The "Infinite Retry" Loop (Why we need Idempotency)

During a broker API outage, an agent correctly identified that a trade failed. However, the agent's recovery logic was "Try again." Without a runtime-managed retry policy, the agent spammed the API, eventually getting the IP banned.

- *Lesson:* Agents should not handle retries. The Runtime handles reliability using **Idempotency Keys** and **Exponential Backoff**.

11.3 Debugging the "Black Box" (Why we need DFID)

After a week of operation, the portfolio balance drifted inexplicably. Logs showed *what* trades happened, but not *why*. I spent days grepping through scattered text logs to find the prompt that caused a specific bad trade.

- *Lesson:* Logs are useless without correlation. Implementing **DecisionFlow ID (DFID)** allowed me to visualize the entire chain: *Market Signal -> Prompt -> Agent Thought -> Policy -> Action*.

11.4 What's Next

This document outlines the architectural pattern. In a forthcoming article, I will detail the specific implementation of **Alvestor**, demonstrating how these abstract concepts map to a concrete stack (Python, Event Bus, SQL). Additionally, I plan to release a set of **ROA/DIR Code Templates** in this repository to help developers bootstrap reliable agents without reinventing the wheel.

12. Observability & Metrics

To pass a "Production Readiness" audit, the Runtime must be observable. Operators need to know not just *what* happened, but *how* the system is performing. DIR mandates the export of the following golden signals:

Metric	Type	Description
validation_latency_ms	Histogram	Time spent in DIM. Split by <code>type=hard</code> (code) and <code>type=soft</code> (LLM).
stale_context_reject_rate	Counter	Number of proposals rejected by JIT State Check (<code>STATE_DRIFT</code>). High rates indicate need for optimization.
escalation_count	Counter	Total escalations, tagged by <code>agent_id</code> and <code>reason</code> . Helps identify "needy" agents.
resource_lock_contention	Gauge	Number of active locks vs. configured limits.
token_budget_burn	Histogram	Tokens consumed per DecisionFlow.
idempotency_hit_rate	Counter	Number of times the Runtime saved a redundant API call.

These metrics should be scraped by a standard observability stack (Prometheus/Grafana) to visualize the "health" of the cognitive architecture.

13. Trade-offs and Limitations

Engineering is about trade-offs. Adopting DIR introduces friction and cost. It is not a silver bullet.

12.1 Latency vs. Safety

DIR introduces a "Validation Tax." Every decision must be serialized, validated, and logged before execution.

- *Trade-off:* In High-Frequency Trading (HFT), where microseconds matter, DIR is too slow.
- *Fit:* DIR is optimized for "Human-Speed" or "Business-Speed" decisions (seconds to minutes), where safety outweighs raw speed.

Why Blocking Validation Matters

In many recommender systems, bad decisions affect "reputation" or "ranking" asynchronously. In financial or physical systems, a bad decision creates an irreversible loss (financial or kinetic). Therefore, DIR favors **Pre-flight Validation** over **Post-flight Audit**. While this introduces latency, it enforces the "Safety First" invariant. A missed trade due to aggressive validation is an opportunity cost; a realized hallucination is an actual loss. The system is designed to fail safe (reject) rather than fail open (execute and apologize).

12.2 Complexity Overhead

Implementing a Context Compiler, Policy Engine, and State Machine is significantly harder than writing a `while True` loop with an LLM.

- *Trade-off:* For simple experimental scripts, DIR is over-engineering. It pays off only when the cost of failure is non-zero (e.g., handling money, PII data, or customer interactions).

12.3 The "Human Bottleneck" in Policy Design

DIR moves the complexity from "Prompt Engineering" to "Policy Engineering." Defining the JSON schemas, RBAC roles, and validation rules (Rego/Python) requires domain expertise that cannot be automated. You still need humans to define *what* is safe.

12.4 The "Garbage Policy In" Risk

DIR guarantees that an agent cannot violate the *syntax* or *permissions* of the system. However, it cannot guarantee that a syntactically valid decision is *smart*. If a policy allows an agent to "Delete Database" and the agent requests it in a valid JSON format, the Runtime will execute it. **Safety is only as strong as the weakest rule in the Policy Enforcement Point.** This places a burden on "Policy Engineering"-developers must define granular, least-privilege constraints (e.g., "Delete only records created by this agent"). DIR acts as a "Safety Kernel," not an "Oracle"; it guarantees that an action is permissible, not that it is wise. Safety is a property of the whole system, not just the model.

12.5 Adversarial Robustness (Defense in Depth)

While DIR prevents accidents, it is not a silver bullet against targeted adversarial attacks. If a malicious actor successfully performs a specific prompt injection that forces the LLM to output a valid, authorized, but malicious Policy Proposal (e.g., "Sell at minimum allowable price"), limits may be respected but intent subverted. DIR is a layer of **Defense in Depth**^[12]; it must be complemented by upstream defenses like input sanitization and semantic monitoring.

12.6 Scope of the OS Analogy

The Kernel Space / User Space analogy used throughout DIR is **architectural, not mechanistic**. LLM inference does not work like a system call. Token sampling is not equivalent to instruction execution. The analogy is pedagogical, not literal.

What the analogy *does* capture accurately:

- **Privilege separation:** Unprivileged processes (agents) request actions; the privileged kernel validates and executes them.
- **Failure isolation:** A crashing user process cannot corrupt kernel state — a hallucinating agent cannot corrupt authoritative system state.
- **Access control:** User processes cannot self-grant permissions — agents cannot self-grant execution authority.

What the analogy does *not* imply:

- That DIR reduces the *probability* of agent failure — it does not. Agents will hallucinate regardless of runtime discipline.
- That the mechanisms of OS protection rings map to LLM internals — they do not.
- That deterministic validation eliminates probabilistic risk — it contains its *consequences*, not its *occurrence*.

The goal of DIR is containment, not prevention. An agent can be wrong, loop, or hallucinate in User Space — the Kernel prevents those failures from propagating into irreversible side effects. This is analogous to how an OS kernel protects filesystem integrity from a buggy process, without making the process less buggy.

A similar separation — arriving independently from delegation theory rather than systems engineering — is documented in Google DeepMind's *Intelligent AI Delegation* (arXiv:2602.11865, 2026), which converges on the same primitive: decoupling reasoning authority from execution authority.

13. Conclusion: From Chatbots to Systems

We are currently in the "wild west" phase of Agentic AI. Developers are connecting powerful probabilistic models directly to sensitive APIs, relying on "prompt injection defense" as their only safety net.

This is unsustainable.

As we move from building chatbots to building autonomous systems, we must stop treating AI as a magic box and start treating it as an untrusted component in a critical system.

The Decision Intelligence Runtime (DIR) is an architectural response to this reality. By separating **Reasoning** (the Agent) from **Execution** (the Runtime), and enforcing strict invariants like **Idempotency**, **Temporal Validity**, and **Auditability**, we can build systems that are not just "smart," but also safe, reliable, and accountable.

Alvestor proved that an LLM can trade stocks without going broke—but only when it is put in a straightjacket of deterministic engineering.

14. Glossary

- **Decision Intelligence Runtime (DIR):** An architectural pattern that separates probabilistic agent reasoning from deterministic execution to ensure safety and reliability.
- **Responsibility-Oriented Agents (ROA):** A design pattern for bounding AI autonomy where agents operate within strict functional contracts.
- **DecisionFlow ID (DFID):** A unique correlation identifier that traces the entire lifecycle of a decision, from trigger to execution result.
- **Decision Validity Window (DVW):** The time period during which a proposed decision is considered valid. If execution is attempted after this window, it is rejected.
- **Decision Integrity Module (DIM):** The deterministic component of the runtime acting as a gatekeeper (Policy Enforcement Point) that validates all proposals.
- **Policy Proposal:** A structure defining the agent's desired action (intent) submitted to the runtime for validation.
- **Execution Intent:** A validated and approved object derived from a Policy Proposal, authorized to trigger external side effects.
- **Policy Enforcement Point (PEP):** A security architectural component that acts as a gatekeeper, intercepting requests and validating them against a set of policies before allowing execution. In DIR, the Decision Integrity Module (DIM) functions as the PEP.
- **Context Compilation:** The process of deterministically assembling a snapshot of relevant data (state, history, rules) for the agent before invocation.
- **Semantic Alignment Check:** A validation step that compares the agent's natural language explanation with its structured policy to detect hallucinations or mismatches.
- **Escalation Budget:** A rate-limiting mechanism that restricts how many times an agent can request human intervention within a given timeframe.

15. References

1. **Policy Enforcement Point (PEP)** - Defined in [RFC 2753](#) and adopted in [NIST SP 800-207 Zero Trust Architecture](#). ↩
2. **Open Policy Agent (OPA)** - The standard for policy-as-code. [openpolicyagent.org](#). ↩
3. **Service Discovery** - Richardson, C. Microservices Patterns. [microservices.io/patterns/client-side-discovery.html](#). ↩
4. **CQRS (Command Query Responsibility Segregation)** - Fowler, M. [martinfowler.com/bliki/CQRS.html](#). ↩
5. **Distributed Tracing** - The underlying concept for OpenTelemetry. [opentelemetry.io/docs/concepts/signals/traces/](#). ↩
6. **Declarative API** - A core principle of Kubernetes. [kubernetes.io/docs/concepts/overview/working-with-objects/kubernetes-objects/](#). ↩
7. **Optimistic Concurrency Control** - Fowler, M. [martinfowler.com/eaCatalog/optimisticOfflineLock.html](#). ↩
8. **Idempotency** - Ensuring requests can be retried safely. See [Stripe API Idempotency](#) or [IETF Draft](#). ↩
9. **Saga Pattern** - Managing failures in distributed transactions. [microservices.io/patterns/data/saga.html](#). ↩
10. **Retrieval-Augmented Generation (RAG)** - Lewis et al. (2020). [arxiv.org/abs/2005.11401](#). ↩
11. **Event-Driven Architecture** - Decoupling services via events. [martinfowler.com/articles/201701-event-driven.html](#). ↩
12. **Defense in Depth** - A security strategy overlapping multiple defensive layers. [csrc.nist.gov/glossary/term/defense_in_depth](#). ↩